

5 Flower Types Classification Dataset

The **5 Flower Types Classification Dataset** is a collection of images belonging to five different flower classes: Lilly, Lotus, Sunflower, Orchid, and Tulip. Each flower class contains 1000 images, resulting in a total of 5000 images in the dataset.

This dataset is suitable for training and evaluating a multi-class Convolutional Neural Network (CNN) model to classify flower images into one of the five mentioned classes. The goal of the classification task is to accurately identify the type of flower from an input image.

The dataset can be used to explore various deep learning techniques for image classification, such as data augmentation, transfer learning, and model fine-tuning. It provides a challenging task due to the visual similarity and subtle differences among different flower types.

Dataset Details:

- Number of classes: 5
- Total images: 5000 (1000 images per class)
- Image format: JPG or PNG
- Image resolution: Varies (please preprocess the images to a consistent size if required)

The 5 Flower Types Classification Dataset is a valuable resource for researchers, students, and practitioners interested in the field of computer vision, specifically in image classification tasks. It can be used for educational purposes, benchmarking different models, and advancing the state-of-the-art in flower classification.

Feel free to download the dataset and start exploring the fascinating world of flower image classification!:

<https://www.kaggle.com/datasets/kausthubkannan/5-flower-types-classification-dataset>

(<https://www.kaggle.com/datasets/kausthubkannan/5-flower-types-classification-dataset>)

```
In [1]: 1 # import libraries
        2 import os
        3 import shutil
        4 import random
        5 import numpy as np
        6 import pandas as pd
        7 from matplotlib import pyplot as plt
        8 from PIL import Image
        9 from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
       10 import tensorflow as tf
       11 from tensorflow.keras.preprocessing.image import ImageDataGenerator
       12 from tensorflow.keras.optimizers import Adam
       13 from tensorflow.keras.layers import Conv2D, MaxPool2D, Flatten, Dense, Dropout
       14 from tensorflow.keras import Sequential
       15 os.environ['KAGGLE_CONFIG_DIR'] = '/content'
```

Downloading and preparing data for model

```
In [2]: 1 !kaggle datasets download -d kausthubkannan/5-flower-types-classification-dataset
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /content/kaggle.json'

Downloading 5-flower-types-classification-dataset.zip to /content

96% 232M/242M [00:02<00:00, 104MB/s]

100% 242M/242M [00:02<00:00, 115MB/s]

```
In [3]: 1 !unzip \*.zip && rm *.zip
```

Streaming output truncated to the last 5000 lines.

```
inflating: flower_images/Lilly/00048a5c76.jpg
inflating: flower_images/Lilly/001ff6644e.jpg
inflating: flower_images/Lilly/001ff6656j.jpg
inflating: flower_images/Lilly/00973ad1b1.jpg
inflating: flower_images/Lilly/00a7d512d6.jpg
inflating: flower_images/Lilly/00f36a3c40.jpg
inflating: flower_images/Lilly/013628cccc.jpg
inflating: flower_images/Lilly/01998d6fb5.jpg
inflating: flower_images/Lilly/01a0ec319c.jpg
inflating: flower_images/Lilly/01b4bb0289.jpg
inflating: flower_images/Lilly/025ef3ea44.jpg
inflating: flower_images/Lilly/02a7a2df46.jpg
inflating: flower_images/Lilly/02be2ca388.jpg
inflating: flower_images/Lilly/035cce082f.jpg
inflating: flower_images/Lilly/039eba79d4.jpg
inflating: flower_images/Lilly/04067b91d6.jpg
inflating: flower_images/Lilly/04acfd5449.jpg
inflating: flower_images/Lilly/05777790e2.jpg
```

```
In [4]: 1 !pip install split-folders
```

Looking in indexes: <https://pypi.org/simple>, (<https://pypi.org/simple>,) <https://us-python.pkg.dev/colab-wheels/public/simple/> (<https://us-python.pkg.dev/colab-wheels/public/simple/>)
Collecting split-folders
 Downloading split_folders-0.5.1-py3-none-any.whl (8.4 kB)
Installing collected packages: split-folders
Successfully installed split-folders-0.5.1

```
In [5]: 1 import splitfolders
```

```
In [6]: 1 src_dir = '/content/flower_images'
2 dst_dir = '/content/Data'
```

Data preprocessing (image augmentation)

```
In [7]: 1 splitfolders.ratio(input=src_dir, output=dst_dir, ratio=(0.8, 0.2))
```







Copying files: 5000 files [00:00, 5625.05 files/s]

```
In [8]: 1 train_datagen = ImageDataGenerator(rescale=1/255., rotation_range=0.2,
2                                     # brightness_range=(0.2, 0.5),
3                                     zoom_range=0.2, shear_range=0.2,
4                                     horizontal_flip=True)
5 train_dataset = train_datagen.flow_from_directory('/content/Data/train',
6                                     target_size=(150, 150),
7                                     batch_size=32,
8                                     shuffle=True)
9
10
11 val_datagen = ImageDataGenerator(rescale=1/255.)
12 val_dataset = val_datagen.flow_from_directory('/content/Data/val', target_size=(150, 150),
13                                     batch_size=32, shuffle=False)
```

Found 4000 images belonging to 5 classes.
Found 1000 images belonging to 5 classes.

In [9]:

```
1 images, labels = next(train_dataset)
2 labels = np.argmax(labels, axis=1)
3 class_names = list(train_dataset.class_indices.keys())
4 def plot_random_images(images, labels, class_names):
5     plt.figure(figsize=(12, 6))
6
7     for i in range(6):
8         ax = plt.subplot(2, 3, i+1)
9         rand_index = random.choice(range(len(images)))
10        plt.imshow(images[rand_index])
11        plt.title(class_names[labels[rand_index]], color='green', fontsize=12)
12
13    plt.tight_layout()
14    plt.show()
15
16 plot_random_images(images, labels, class_names)
```



creating model

```
In [10]: 1 model = Sequential([
2
3           Conv2D(filters=16, kernel_size=(3,3), strides=1, activation='relu', i
4           MaxPool2D(pool_size=(2,2), strides=2, padding='valid'),
5
6           Conv2D(filters=32, kernel_size=(3,3), strides=2, activation='relu'),
7           MaxPool2D(pool_size=(2,2), strides=1, padding='same'),
8
9           Conv2D(filters=64, kernel_size=(3,3), strides=2, activation='relu'),
10          MaxPool2D(pool_size=(2,2), strides=1, padding='same'),
11
12          Flatten(),
13          Dense(256, activation='relu'),
14          Dense(128, activation='relu'),
15          Dense(64, activation='relu'),
16          Dense(5, activation='softmax')
17
18      ])
```

```
In [11]: 1 model.compile(optimizer=Adam(), loss=tf.keras.losses.CategoricalCrossentropy(), metrics=[
```

```
In [12]: 1 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 148, 148, 16)	448
max_pooling2d (MaxPooling2D)	(None, 74, 74, 16)	0
conv2d_1 (Conv2D)	(None, 36, 36, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 32)	0
conv2d_2 (Conv2D)	(None, 17, 17, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 64)	0
flatten (Flatten)	(None, 18496)	0
dense (Dense)	(None, 256)	4735232
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 5)	325
=====		
Total params: 4,800,293		
Trainable params: 4,800,293		
Non-trainable params: 0		

In [13]:

```
1 history = model.fit(train_dataset, epochs=15, validation_data=(val_dataset))
```

Epoch 1/15

125/125 [=====] - 69s 411ms/step - loss: 1.2671 - accuracy: 0.4565
- val_loss: 1.1449 - val_accuracy: 0.5250

Epoch 2/15

125/125 [=====] - 50s 399ms/step - loss: 1.0477 - accuracy: 0.5885
- val_loss: 1.0282 - val_accuracy: 0.6100

Epoch 3/15

125/125 [=====] - 48s 387ms/step - loss: 0.9237 - accuracy: 0.6470
- val_loss: 0.9004 - val_accuracy: 0.6420

Epoch 4/15

125/125 [=====] - 48s 386ms/step - loss: 0.8279 - accuracy: 0.6775
- val_loss: 0.8699 - val_accuracy: 0.6850

Epoch 5/15

125/125 [=====] - 49s 390ms/step - loss: 0.7596 - accuracy: 0.7138
- val_loss: 0.8166 - val_accuracy: 0.6950

Epoch 6/15

125/125 [=====] - 50s 398ms/step - loss: 0.6770 - accuracy: 0.7347
- val_loss: 0.7740 - val_accuracy: 0.7160

Epoch 7/15

125/125 [=====] - 53s 422ms/step - loss: 0.5846 - accuracy: 0.7760
- val_loss: 0.7458 - val_accuracy: 0.7430

Epoch 8/15

125/125 [=====] - 49s 390ms/step - loss: 0.5382 - accuracy: 0.8023
- val_loss: 0.7037 - val_accuracy: 0.7670

Epoch 9/15

125/125 [=====] - 48s 387ms/step - loss: 0.4564 - accuracy: 0.8330
- val_loss: 0.6574 - val_accuracy: 0.7850

Epoch 10/15

125/125 [=====] - 49s 395ms/step - loss: 0.3924 - accuracy: 0.8593
- val_loss: 0.7169 - val_accuracy: 0.7680

Epoch 11/15

125/125 [=====] - 52s 418ms/step - loss: 0.3647 - accuracy: 0.8748
- val_loss: 0.7141 - val_accuracy: 0.7810

Epoch 12/15

125/125 [=====] - 48s 388ms/step - loss: 0.3328 - accuracy: 0.8815
- val_loss: 0.6606 - val_accuracy: 0.8000

Epoch 13/15

125/125 [=====] - 49s 391ms/step - loss: 0.2856 - accuracy: 0.9035
- val_loss: 0.6334 - val_accuracy: 0.8220

Epoch 14/15

125/125 [=====] - 50s 399ms/step - loss: 0.2379 - accuracy: 0.9160
- val_loss: 0.7007 - val_accuracy: 0.8260

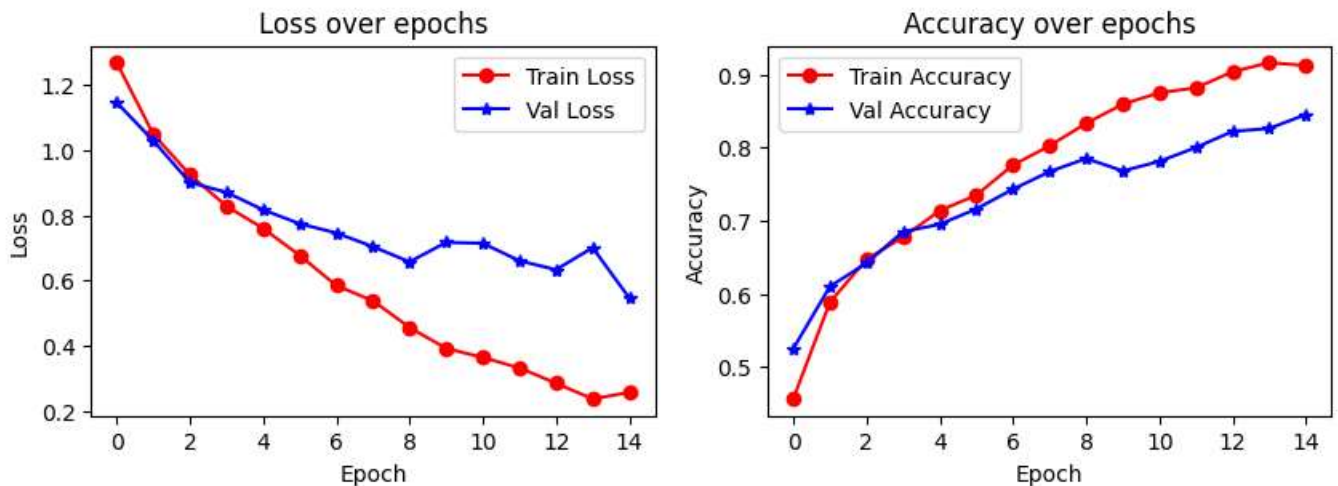
Epoch 15/15

125/125 [=====] - 48s 389ms/step - loss: 0.2576 - accuracy: 0.9120
- val_loss: 0.5470 - val_accuracy: 0.8450

```

In [14]: 1 loss_df = pd.DataFrame(history.history)
2
3 def plot_predictions(data=loss_df):
4     fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(10, 3))
5
6     ax1.plot(loss_df['loss'], color='red',marker='o', label='Train Loss')
7     ax1.plot(loss_df['val_loss'], color='blue',marker='*', label='Val Loss')
8
9     ax1.set_title('Loss over epochs')
10    ax1.set_xlabel('Epoch')
11    ax1.set_ylabel('Loss')
12    ax1.legend()
13
14    ax2.plot(loss_df['accuracy'], color='red',marker='o', label='Train Accuracy')
15    ax2.plot(loss_df['val_accuracy'], color='blue',marker='*', label='Val Accuracy')
16
17    ax2.set_title('Accuracy over epochs')
18    ax2.set_xlabel('Epoch')
19    ax2.set_ylabel('Accuracy')
20    ax2.legend()
21 plot_predictions(loss_df)

```



making predictions

```

In [15]: 1 predictions = np.argmax(model.predict(val_dataset), axis=1)
2         y_true = val_dataset.labels

```

32/32 [=====] - 6s 172ms/step

```


In [16]: 1 print(accuracy_score(predictions, y_true))

```

0.845

In [17]:

```
1 def plot_confusion_matrix(y_true, predictions, class_names):
2
3     cm = confusion_matrix(y_true, predictions)
4     plt.figure(figsize=(8, 6))
5     heatmap = plt.imshow(cm, cmap='Blues')
6
7     # Set axis labels and title
8     plt.xlabel('Predicted Labels')
9     plt.ylabel('True Labels')
10    plt.title('Confusion Matrix')
11
12    # Set xticks and yticks with class names
13    tick_labels = class_names
14    plt.xticks(ticks=np.arange(len(class_names)), labels=tick_labels, rotation=45)
15    plt.yticks(ticks=np.arange(len(class_names)), labels=tick_labels)
16
17    # Add numbers to the heatmap cells
18    for i in range(len(class_names)):
19        for j in range(len(class_names)):
20            plt.text(j, i, str(cm[i, j]), ha='center', va='center', color='black')
21
22    plt.colorbar(heatmap)
23    plt.show()
24    plot_confusion_matrix(y_true, predictions, class_names)
```



In [20]:

```
1 def plot_random_image(model, val_data, classes):
2
3     images = []
4     labels = []
5     for _ in range(len(val_data)):
6         batch_images, batch_labels = next(val_data)
7         images.extend(batch_images)
8         labels.extend(batch_labels)
9
10    # Shuffle the images and labels together
11    combined = list(zip(images, labels))
12    random.shuffle(combined)
13    images, labels = zip(*combined)
14    labels = np.argmax(labels, axis=1)
15    plt.figure(figsize=(12, 6))
16    for i in range(6):
17        ax = plt.subplot(2, 3, i + 1)
18        rand_index = random.choice(range(len(images)))
19        target_image = images[rand_index]
20        pred_probs = model.predict(tf.expand_dims(target_image, axis=0), verbose=0)
21        pred_label = classes[pred_probs.argmax()]
22        true_label = classes[labels[rand_index]]
23
24        plt.imshow(target_image)
25
26        if pred_label == true_label:
27            color = "green"
28        else:
29            color = "red"
30
31        plt.title("Pred: {} {:.20f}% (True: {})".format(pred_label,
32                                                       100 * tf.reduce_max(pred_probs),
33                                                       true_label),
34                color=color, fontsize=10)
35
36    plt.tight_layout()
37    plot_random_image(model, val_dataset, class_names)
38    plt.show()
```

